

Prediction for Economists

Prediction as an objective


- A rule: observed characteristics in, an unobserved outcome out.
- Estimated on one sample, judged on another.
- The criterion is error on the second, not the value of any coefficient.
- Bias in the coefficients is acceptable when it lowers that error.

Everything that follows is about estimating that error honestly.

Outline

1. Overfitting and the trade-off
2. Estimating out-of-sample error
3. Leakage
4. Measures of accuracy
5. Constructing regressors
6. Choosing an estimator
7. Putting it together


Overfitting



- In-sample fit: better at every degree.
- 20 observations, degree 19: every point exactly.
- Error in an independent sample: lowest at a low degree.


Fit in the estimation sample is not evidence of accuracy elsewhere.

Overfitting with a binary outcome



- Overfitting is not a property of polynomials or of continuous outcomes.
- Left unrestricted, the tree keeps splitting until each region is pure.
- It ends up isolating single observations, and misclassifies new ones.


Bias and variance



- Each grey line: one independent sample of 25 observations.
- Restrictive fit: stable across samples, wrong on average.
- Flexible fit: right on average, far off in any one sample.

Two distinct errors, and less of one buys more of the other.

The bias–variance trade-off



- Squared bias, variance, noise.
- Restrictive: bias dominates.
- Flexible: variance dominates.
- The noise is irreducible.
- One point on the curve is best.

Which point, though? The curve is drawn here because the truth is known.

Outline

1. Overfitting and the trade-off
2. Estimating out-of-sample error
3. Leakage
4. Measures of accuracy
5. Constructing regressors
6. Choosing an estimator
7. Putting it together

Holding observations back

Training	Validation	Test
----------	------------	------

- Out-of-sample error is unobserved. Hide observations and imitate it.
- Training: estimate the coefficients of one given specification.
- Validation: measure error there to *choose* — which regressors, which estimator, how flexible.
- Test: measure error once more, to *report* what the chosen rule will do.


Both measure error; they answer different questions.

Why the test sample is separate

- Compare twenty specifications on the validation sample. One wins.
- It won partly on merit, partly because that sample's noise suited it.
- So the winner's validation error is better than its true accuracy.
- The more specifications compared, the larger that gap.

The test sample took part in no comparison, so its error is not flattered by having been chosen on. That is the number to report.

Cross-validation



- One validation block wastes those observations, and is one arbitrary draw.
- Instead, cut the training data into k folds of equal size.
- Each fold takes the validation role in turn; the other $k - 1$ estimate.
- Average the k errors: one number per specification, using all the data.

The validation role rotates inside the training data; nothing is permanently set aside. The test sample stays outside the rotation, still unused.

Outline


1. Overfitting and the trade-off
2. Estimating out-of-sample error
3. Leakage
4. Measures of accuracy
5. Constructing regressors
6. Choosing an estimator
7. Putting it together

What leakage is

- Leakage: the fit uses information that will not be available when the rule is actually applied.
- The estimate of accuracy is then an estimate of a task that is easier than the real one.
- It shows up as a good cross-validated error and a disappointing one in use.

Nothing in the procedure complains. The fold was held out, so the arithmetic is correct; it is the data that were contaminated.

Leakage in a cross-validated fit



- Here one regressor was built from the outcome itself.
- Cross-validation reports an almost perfect fit.
- Remove that regressor and the fit is ordinary.

The held-out fold was held out, but its outcome had already travelled into the regressors.

Two kinds of leakage

- Wrong date: a regressor recorded after the outcome, or unknown at the moment the prediction would be made.
- Wrong sample: a quantity computed on the whole data and then used inside the folds.
- The second is the common one: standardization, imputation, category means, the choice of which regressors to keep.

Both make the cross-validated error too low, and the error in use a surprise.

Keeping the estimate honest


- Split first, transform afterwards. Never the reverse.
- Means, standard deviations and imputed values are estimated on the training fold, then applied to the held-out one.
- A regressor enters only if it is known before the outcome is realized.
- The test sample is opened once, at the end.

In code, this is why the transformations and the estimator are fitted together, inside the fold, rather than beforehand.

Outline

1. Overfitting and the trade-off
2. Estimating out-of-sample error
3. Leakage
4. Measures of accuracy
5. Constructing regressors
6. Choosing an estimator
7. Putting it together



Measuring accuracy: continuous outcomes



- Root mean squared error: large errors count for more.
- Mean absolute error: every unit of error counts the same.
- A few extreme observations move the first, not the second.

Fix the measure before the estimation, by the cost of an error.

Measuring accuracy: binary outcomes




- Rare ones: the share of correct classifications says little.
- Zero everywhere: 96% correct, nothing identified.
- Predicted against actual, as a table: the useful summary.

Thresholds, precision, and recall

- Estimators return a probability. A classification needs a threshold.
- Precision: of those predicted one, the share that are one.
- Recall: of those that are one, the share predicted one.
- A higher threshold: more precision, less recall.

The threshold is a decision about the cost of the two errors, not an estimation problem.

The ROC curve




- Both rates, at every threshold.
- Its area: the quality of the ordering.
- No one threshold involved.
- The diagonal: no information.

Outline

1. Overfitting and the trade-off
2. Estimating out-of-sample error
3. Leakage
4. Measures of accuracy
5. Constructing regressors
6. Choosing an estimator
7. Putting it together

Lesson 1: an estimator adds or splits; it never combines



- Both estimators hold the ingredients, and fifteen irrelevant regressors.
- Neither builds the combination. Both are transformed by receiving it.
- In the second panel the error falls by about 60% for each of them.

Lesson 1, stated exactly

- A linear model is additive in the regressors it is handed. It can weight them; it cannot multiply or divide them.
- A tree is a step function on rectangles. A ratio is not a step function, so no finite tree equals one.
- A group mean, a rank within a group, a lag: each needs an operation across rows, and neither family performs operations across rows.

This is why the combination has to be written down. It is not hidden in the data waiting to be found; it is not in the space of functions either estimator can represent.

Which combinations

- Ratios: debt to assets, price to earnings, output per worker.
- Deflated, real and per-capita terms, so units compare across time and size.
- Differences and growth rates: the level and the change are different regressors.
- Position within a group: size against the industry median, income against the region, a rank within the cohort.


An economist can write these down in a morning. No estimator will discover them in any amount of data.

Lesson 2: encoding a category by its mean outcome

- A category with hundreds of levels — municipality, product code, four-digit industry — gives hundreds of indicators, each carried by a handful of observations.
- Target encoding replaces the level by one number: the mean outcome among the observations in it. Madrid becomes 0.34.
- Rare levels are shrunk towards the overall mean, $(n_\ell \bar{y}_\ell + m \bar{y}) / (n_\ell + m)$.
- Hundreds of regressors become one, and it feeds any estimator.

The regressor is built out of the outcome, which is what makes it strong and what makes it dangerous.

Lesson 2: when the means are computed decides everything



- 400 levels, five observations each, independent of the outcome.
- Whole sample: every row sits inside its own regressor, weight a fifth.
- Training fold: the training rows still do, and a false slope is learned.

Zero is the honest answer, and only the third column gives it. Each row needs a mean computed without it.

Others, briefly

- Permutation importance: shuffle a regressor, see how much the cross-validated error worsens. Any estimator.
- Principal components: many correlated regressors replaced by a few, chosen without reference to the outcome.
- Lags, differences and rolling means, where the data have a time dimension.
- The outcome can be transformed too: squared error in logs penalizes relative error.

Permutation importance doubles as a leakage check. One regressor carrying nearly all of the accuracy is the signature.

Outline

1. Overfitting and the trade-off
2. Estimating out-of-sample error
3. Leakage
4. Measures of accuracy
5. Constructing regressors
6. Choosing an estimator
7. Putting it together

Three decisions

- Which family: penalized linear regression, or something tree-based.
- How flexible within that family: the penalty, the depth, the number of trees.
- When to stop: a gain smaller than the variation across folds is not a gain.

All three are settled by the same number, the cross-validated error, computed on identical folds with identical regressors. Least squares is the first entry, as the benchmark the rest have to beat.

The first decision: which family

Restrictive fits win	Flexible fits win
Few observations per regressor	Many observations
Weak signal, much noise	Strong signal
Smooth, additive relations	Thresholds and interactions
Extrapolation is required	Prediction inside the data


Most outcomes in economics sit on the left: many regressors, modest signal. Each row is a conjecture, and the folds settle it.

The second decision: what each estimator asks you to choose

- Least squares: nothing.
- Ridge and LASSO: one penalty. LASSO also sets coefficients to zero.
- Tree: the depth, or the smallest region allowed.
- Random forest: the number of trees and the regressors per split. Little turns on either.
- Boosting: the rate, the depth, and the number of trees. Much turns on all three.

Each estimator has its counterpart for a binary outcome, with the same choices to make.


Penalized regression



- 70 observations, 40 regressors, 3 with nonzero coefficients.
- The penalty shrinks the coefficients, and sets the irrelevant to zero.
- Lowest test error at a strictly positive penalty.
- Which regressors survive is not a finding: another sample selects a different set and predicts about as well.

The penalty is one more choice, and the folds make it.

Trees and random forests




- One tree: the boundary is a feature of the particular sample.
- A forest: many trees on resampled data, averaged.
- The flexibility kept, most of the variance gone.

Boosting

- Trees in sequence, not in parallel.
- Each one fitted to the residuals of the current model.
- Each entering with a small weight, so progress is gradual.
- More parameters than a forest, and it overfits when they are careless.

A forest asks fewer questions of you, and is the better place to start.

Comparing estimators




- Same regressors, same folds, same measure, for every estimator.
- Bars: the average over the five folds. Whiskers: one standard error.
- Flexible fits win here; on data close to linear the ordering reverses.

The third decision, read off the whiskers: on the left boosting is genuinely ahead of the forest, on the right the two are within one standard error and the forest is the simpler rule.

Outline

1. Overfitting and the trade-off
2. Estimating out-of-sample error
3. Leakage
4. Measures of accuracy
5. Constructing regressors
6. Choosing an estimator
7. Putting it together

The workflow



- Was every transformation fitted inside the fold?
- Could any regressor contain the outcome?
- Same regressors and same folds across every specification?

Step 6 happens once. Anything learned there is not a reason to return to step 2.

Summary

- Accuracy concerns observations not used in the estimation.
- In-sample fit does not measure it; bias against variance is the trade-off.
- Cross-validation estimates it. Leakage quietly destroys the estimate.
- Fix the measure of accuracy in advance, by the cost of an error.
- Regressors first, estimator second.

Appendix


The decomposition of expected prediction error

For a new observation at x_0 , with $y = f(x_0) + \varepsilon$ and $\hat{f}$ estimated on a random sample,

$$\mathbb{E}[(y - \hat{f}(x_0))^2] = (\mathbb{E}[\hat{f}(x_0)] - f(x_0))^2 + \text{Var}(\hat{f}(x_0)) + \text{Var}(\varepsilon).$$

- The expectation is taken over samples, not over observations.
- The first two terms can be traded against one another.
- The third cannot be reduced by any estimator.
- Penalization and averaging lower the second at some cost in the first.


Sampling variability of a single split



- Each point: a new partition.
- One split varies a great deal.
- The average over folds is stable.
- Small samples: it matters most.

With a binary outcome, keep the share of ones equal across folds.

Boosting: successive fits




- A single tree of depth two is a coarse step function.
- Successive trees fitted to the residuals accumulate into an accurate fit.
- Continued far enough, the procedure begins to fit the noise.
- The number of trees is itself chosen by cross-validation.

Bet on sparsity

- If many regressors each matter a little, no procedure predicts well: the data cannot support that many estimated effects.
- If a few matter a lot, penalized methods find them.
- So act as though few matter. Build the candidates, let the penalty select.
- The bet is not always right. It is that there is nothing better to bet on.

Hastie, Tibshirani and Friedman, sec. 16.2.2. Stated there as a principle, not a theorem.

Transformations of a single regressor



- A tree splits on $x \leq t$, so it reads the ordering of x and nothing else.
- An increasing transformation leaves its partition, and its fit, unchanged.
- The forest scores 0.30 either way; the same log is worth 0.15 to least squares.

Rescaling one regressor: everything to a linear model, exactly nothing to a tree.

Combining two regressors: worth doing for both.

Categorical variables: other encodings

- Few levels: one indicator each, and nothing further is needed.
- Many levels: as many regressors as levels, most of them nearly empty, and each estimated on a handful of observations.
- Group the rare levels together, or replace a level by its frequency.
- The frequency uses no outcome, so it can be computed once, before the split.

Missing values carry information

- Values are rarely missing at random: no credit history, a firm too young to report, a question refused.
- Impute the value, and add an indicator for having imputed it.
- The indicator is often the better predictor of the two.
- Both the imputation and the indicator are built inside the fold.

Dropping the incomplete observations throws that information away, and leaves a sample selected on the regressors.

Further measures for binary outcomes

- When ones are rare, the precision–recall curve is more informative.
- ROC can look favourable when few of the predicted ones are correct.
- The F_1 measure combines precision and recall in one number.
- Log-loss penalizes confident predictions that turn out to be wrong.
- Calibration: of those given probability 0.7, about 70% should be ones.

References

- Mullainathan and Spiess (2017), *Journal of Economic Perspectives* 31(2).
- Athey and Imbens (2019), *Annual Review of Economics* 11, 685–725.
- Domingos (2012), *Communications of the ACM* 55(10), 78–87.
- James, Witten, Hastie and Tibshirani, *Introduction to Statistical Learning*.
- Hastie, Tibshirani and Friedman, *The Elements of Statistical Learning*.